
ELF for the Arm 64-bit Architecture (AArch64)

Release 2020Q2

Document Number: IHI0056G

Arm Ltd

Jun 12, 2020

CONTENTS

1	Preamble	2
1.1	ILP32 Beta	2
1.2	Abstract	2
1.3	Keywords	2
1.4	How to find the latest release of this specification or report a defect in it	2
1.5	Licence	2
1.6	Non-Confidential Proprietary Notice	3
1.7	Contents	3
2	About this document	5
2.1	Change control	5
2.2	References	6
2.3	Terms and abbreviations	6
2.4	Your licence to use this specification	7
3	About This Specification	9
3.1	ELF Class variants	9
4	Platform standards (Example Only)	10
4.1	Linux Platform ABI (example only)	10
5	Object Files	12
5.1	Introduction	12
5.2	ELF Header	12
5.3	Sections	13
5.4	String Table	14
5.5	Symbol Table	14
5.6	Weak Symbols	15
5.7	Relocation	17
6	Program Loading and Dynamic Linking	34
6.1	Program Header	34
6.2	Program Property	34
6.3	Program Loading	35
6.4	Dynamic Linking	35

Document number: IHI 0056G_beta

Release Version: 2020Q2

PREAMBLE

1.1 ILP32 Beta

This document is a beta proposal for ILP32 extensions to ELF for AArch64. All significant ILP32 changes are highlighted in yellow. Feedback welcome through your normal channels.

1.2 Abstract

This document describes the use of the ELF binary file format in the Application Binary Interface (ABI) for the Arm 64-bit architecture.

1.3 Keywords

ELF, AArch64 ELF, ...

1.4 How to find the latest release of this specification or report a defect in it

Please check the Arm Developer site (<https://developer.arm.com/architectures/system-architectures/software-standards/abi>) for a later release if your copy is more than one year old.

Please report defects in this specification to arm.eabi@arm.com.

1.5 Licence

THE TERMS OF YOUR ROYALTY FREE LIMITED LICENCE TO USE THIS ABI SPECIFICATION ARE GIVEN IN *Your licence to use this specification* (page 7) (Arm contract reference LEC-ELA-00081 V2.0). PLEASE READ THEM CAREFULLY.

BY DOWNLOADING OR OTHERWISE USING THIS SPECIFICATION, YOU AGREE TO BE BOUND BY ALL OF ITS TERMS. IF YOU DO NOT AGREE TO THIS, DO NOT DOWNLOAD OR USE THIS SPECIFICATION. THIS ABI SPECIFICATION IS PROVIDED “AS IS” WITH NO WARRANTIES (SEE *Your licence to use this specification* (page 7) FOR DETAILS).

1.6 Non-Confidential Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

If any of the provisions contained in these terms conflict with any of the provisions of any click through or signed written agreement covering this document with Arm, then the click through or signed written agreement prevails over and supersedes the conflicting provisions of these terms. This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <http://www.arm.com/company/policies/trademarks>.

Copyright © 2011-2020 Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.
110 Fulbourn Road, Cambridge, England CB1 9NJ.
LES-PRE-20349

1.7 Contents

Contents

- *ELF for the Arm® 64-bit Architecture (AArch64)*

- *Preamble* (page 2)
 - * *ILP32 Beta* (page 2)
 - * *Abstract* (page 2)
 - * *Keywords* (page 2)
 - * *How to find the latest release of this specification or report a defect in it* (page 2)
 - * *Licence* (page 2)
 - * *Non-Confidential Proprietary Notice* (page 3)
 - * *Contents* (page 3)
- *About this document* (page 5)
 - * *Change control* (page 5)
 - * *References* (page 6)
 - * *Terms and abbreviations* (page 6)
 - * *Your licence to use this specification* (page 7)
- *About This Specification* (page 9)
 - * *ELF Class variants* (page 9)
- *Platform standards (Example Only)* (page 10)
 - * *Linux Platform ABI (example only)* (page 10)
- *Object Files* (page 12)
 - * *Introduction* (page 12)
 - * *ELF Header* (page 12)
 - * *Sections* (page 13)
 - * *String Table* (page 14)
 - * *Symbol Table* (page 14)
 - * *Weak Symbols* (page 15)
 - * *Relocation* (page 17)
- *Program Loading and Dynamic Linking* (page 34)
 - * *Program Header* (page 34)
 - * *Program Property* (page 34)
 - * *Program Loading* (page 35)
 - * *Dynamic Linking* (page 35)

ABOUT THIS DOCUMENT

2.1 Change control

2.1.1 Current status and anticipated changes

The following support level definitions are used by the Arm ABI specifications:

Release Arm considers this specification to have enough implementations, which have received sufficient testing, to verify that it is correct. The details of these criteria are dependent on the scale and complexity of the change over previous versions: small, simple changes might only require one implementation, but more complex changes require multiple independent implementations, which have been rigorously tested for cross-compatibility. Arm anticipates that future changes to this specification will be limited to typographical corrections, clarifications and compatible extensions.

Beta Arm considers this specification to be complete, but existing implementations do not meet the requirements for confidence in its release quality. Arm may need to make incompatible changes if issues emerge from its implementation.

Alpha The content of this specification is a draft, and Arm considers the likelihood of future incompatible changes to be significant.

The ELF32 variant is at “Beta” release quality.

All other content in this document is at the **Release** quality level.

2.1.2 Change history

Issue	Date	By	Change
00bet3	20th December 2011	LS	Beta release
1.0	22nd May 2013	RE	First public release
1.1-beta	6th November 2013	JP	ILP32 Beta
2018Q4	31st December 2018	OS	Typographical changes
2019Q1	29th March 2019	OS	Add Program Property for BTI and PAC. Update MOV[ZK] related relocations.
2019Q2	30th June 2019	LC	Specify ST0_AARCH64_VARIANT_PCS. Update R_<CLS>_TLS_DTPREL and R_<CLS>_TLS_DTPMOD. Clarify GNU_PROPERTY_AARCH64_FEATURE_1_AND.
2019Q4	30th January 2020	TS	Minor layout changes.
2020Q2	1st July 2020	TS	Specify R_<CLS>_PLT32. Correct minus sign not rendering in section <i>Group relocations</i> (page 26). Adjust table widths for readability.

2.2 References

This document refers to, or is referred to by, the following documents.

Ref	External reference or URL	Title
AAELF64	Source for this document	ELF for the Arm 64-bit Architecture (AArch64).
AAPCS64	IHI 0055	Procedure Call Standard for the Arm 64-bit Architecture
Addenda32	IHI 0045	Addenda to, and Errata in, the ABI for the Arm Architecture
LSB ¹	http://www.linuxbase.org/	Linux Standards Base
SCO-ELF ²	http://www.sco.com/developers/gabi/	System V Application Binary Interface – DRAFT
LINUX_ABI ³	https://github.com/hjl-tools/linux-abi/wiki	Linux Extensions to gABI
SYM-VER ⁴	http://people.redhat.com/drepper/symbol-versioning	GNU Symbol Versioning
TLSDESC ⁵	http://www.fsfla.org/~lxoliva/writeups/TLS/paper-lk2006.pdf	TLS Descriptors for Arm. Original proposal document

2.3 Terms and abbreviations

The ABI for the Arm 64-bit Architecture uses the following terms and abbreviations.

A32 The instruction set named Arm in the Armv7 architecture; A32 uses 32-bit fixed-length instructions.

A64 The instruction set available when in AArch64 state.

AAPCS64 Procedure Call Standard for the Arm 64-bit Architecture (AArch64)

AArch32 The 32-bit general-purpose register width state of the Armv8 architecture, broadly compatible with the Armv7-A architecture.

AArch64 The 64-bit general-purpose register width state of the Armv8 architecture.

ABI Application Binary Interface:

1. The specifications to which an executable must conform in order to execute in a specific execution environment. For example, the *Linux ABI for the Arm Architecture*.
2. A particular aspect of the specifications to which independently produced relocatable files must conform in order to be statically linkable and executable. For example, the C++ ABI for the Arm Architecture, *ELF for the Arm Architecture* (page 1), ...

Arm-based ... based on the Arm architecture ...

ELF32 An ELF object file with a class of ELFCLASS32

ELF64 An ELF object file with a class of ELFCLASS64

ILP32 SysV-like data model where int, long int and pointer are 32-bit

LP64 SysV-like data model where int is 32-bit, but long int and pointer are 64-bit.

¹ <http://www.linuxbase.org/>

² <http://www.sco.com/developers/gabi/>

³ <https://github.com/hjl-tools/linux-abi/wiki>

⁴ <http://www.akkadia.org/drepper/symbol-versioning>

⁵ <http://www.fsfla.org/~lxoliva/writeups/TLS/paper-lk2006.pdf>

Q-o-I Quality of Implementation – a quality, behavior, functionality, or mechanism not required by this standard, but which might be provided by systems conforming to it. Q-o-I is often used to describe the tool-chain-specific means by which a standard requirement is met.

T32 The instruction set named Thumb in the Armv7 architecture; T32 uses 16-bit and 32-bit instructions.

Other terms may be defined when first used.

2.4 Your licence to use this specification

IMPORTANT: THIS IS A LEGAL AGREEMENT (“LICENCE”) BETWEEN YOU (AN INDIVIDUAL OR SINGLE ENTITY WHO IS RECEIVING THIS DOCUMENT DIRECTLY FROM ARM LIMITED) (“LICENSEE”) AND ARM LIMITED (“ARM”) FOR THE SPECIFICATION DEFINED IMMEDIATELY BELOW. BY DOWNLOADING OR OTHERWISE USING IT, YOU AGREE TO BE BOUND BY ALL OF THE TERMS OF THIS LICENCE. IF YOU DO NOT AGREE TO THIS, DO NOT DOWNLOAD OR USE THIS SPECIFICATION.

“Specification” means, and is limited to, the version of the specification for the Applications Binary Interface for the Arm Architecture comprised in this document. Notwithstanding the foregoing, “Specification” shall not include (i) the implementation of other published specifications referenced in this Specification; (ii) any enabling technologies that may be necessary to make or use any product or portion thereof that complies with this Specification, but are not themselves expressly set forth in this Specification (e.g. compiler front ends, code generators, back ends, libraries or other compiler, assembler or linker technologies; validation or debug software or hardware; applications, operating system or driver software; RISC architecture; processor microarchitecture); (iii) maskworks and physical layouts of integrated circuit designs; or (iv) RTL or other high level representations of integrated circuit designs.

Use, copying or disclosure by the US Government is subject to the restrictions set out in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.227-7013 or subparagraphs (c)(1) and (2) of the Commercial Computer Software – Restricted Rights at 48 C.F.R. 52.227-19, as applicable.

This Specification is owned by Arm or its licensors and is protected by copyright laws and international copyright treaties as well as other intellectual property laws and treaties. The Specification is licensed not sold.

1. Subject to the provisions of Clauses 2 and 3, Arm hereby grants to LICENSEE, under any intellectual property that is (i) owned or freely licensable by Arm without payment to unaffiliated third parties and (ii) either embodied in the Specification or Necessary to copy or implement an applications binary interface compliant with this Specification, a perpetual, non-exclusive, non-transferable, fully paid, worldwide limited licence (without the right to sublicense) to use and copy this Specification solely for the purpose of developing, having developed, manufacturing, having manufactured, offering to sell, selling, supplying or otherwise distributing products which comply with the Specification.
2. THIS SPECIFICATION IS PROVIDED “AS IS” WITH NO WARRANTIES EXPRESS, IMPLIED OR STATUTORY, INCLUDING BUT NOT LIMITED TO ANY WARRANTY OF SATISFACTORY QUALITY, MERCHANTABILITY, NONINFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE. THE SPECIFICATION MAY INCLUDE ERRORS. Arm RESERVES THE RIGHT TO INCORPORATE MODIFICATIONS TO THE SPECIFICATION IN LATER REVISIONS OF IT, AND TO MAKE IMPROVEMENTS OR CHANGES IN THE SPECIFICATION OR THE PRODUCTS OR TECHNOLOGIES DESCRIBED THEREIN AT ANY TIME.
3. This Licence shall immediately terminate and shall be unavailable to LICENSEE if LICENSEE or any party affiliated to LICENSEE asserts any patents against Arm, Arm affiliates, third parties who have a valid licence from Arm for the Specification, or any customers or distributors of any of them based upon a claim that a LICENSEE (or LICENSEE affiliate) patent is Necessary to implement the Specification. In this Licence; (i) “affiliate” means any entity controlling, controlled by or under common control with a party (in fact or in law, via voting securities, management control or otherwise) and “affiliated” shall be construed accordingly; (ii) “assert” means to allege infringement

in legal or administrative proceedings, or proceedings before any other competent trade, arbitral or international authority; (iii) “Necessary” means with respect to any claims of any patent, those claims which, without the appropriate permission of the patent owner, will be infringed when implementing the Specification because no alternative, commercially reasonable, non-infringing way of implementing the Specification is known; and (iv) English law and the jurisdiction of the English courts shall apply to all aspects of this Licence, its interpretation and enforcement. The total liability of Arm and any of its suppliers and licensors under or in relation to this Licence shall be limited to the greater of the amount actually paid by LICENSEE for the Specification or US\$10.00. The limitations, exclusions and disclaimers in this Licence shall apply to the maximum extent allowed by applicable law.

Arm Contract reference LEC-ELA-00081 V2.0 AB/LS (9 March 2005)

ABOUT THIS SPECIFICATION

This specification provides the processor-specific definitions required by ELF [SCO-ELF⁶] for AArch64-based systems.

The ELF specification is part of the larger Unix System V (SysV) ABI specification where it forms *Object Files* (page 12) and *Program Loading and Dynamic Linking* (page 34). However, the ELF specification can be used in isolation as a generic object and executable format. *Platform standards (Example Only)* (page 10) covers ELF related matters that are platform specific.

Object Files (page 12) and *Program Loading and Dynamic Linking* (page 34) are structured to correspond to chapters 4 and 5 of the ELF specification. Specifically:

- *Object Files* (page 12) covers object files and relocations
- *Program Loading and Dynamic Linking* (page 34) covers program loading and dynamic linking.

3.1 ELF Class variants

Two different pointer sizes are supported by this specification, which result in two very similar but different ELF definitions.

3.1.1 64-bit Pointers, ELF64

- Code and data using 64-bit pointers are contained in an ELF object file with a class of **ELFCLASS64**.
- Referred to as **ELF64** in this specification.
- Pointer-size is **64 bits**.
- Suitable for use by the LP64 variant of [AAPCS64]

3.1.2 32-bit Pointers, ELF32 (Beta)

- Code and data using 32-bit pointers is contained in an ELF object file with a class of **ELFCLASS32**.
- Referred to as **ELF32** in this specification.
- Pointer-size is **32 bits**.
- Suitable for use by the ILP32 variant of [AAPCS64]

Note: Interlinking is not supported between the ELF32 and ELF64 variants.

⁶ <http://www.sco.com/developers/gabi/>

PLATFORM STANDARDS (EXAMPLE ONLY)

We expect that each operating system that adopts components of this ABI specification will specify additional requirements and constraints that must be met by application code in binary form and the code-generation tools that generate such code.

As an example of the kind of issue that must be addressed *Linux Platform ABI (example only)* (page 10), below, lists some of the issues addressed by the *Linux Standard Base* [LSB⁷] specifications.

4.1 Linux Platform ABI (example only)

4.1.1 Symbol Versioning

The Linux ABI uses the GNU-extended Solaris symbol versioning mechanism [SYM-VER⁸].

Concrete data structure descriptions can be found in `/usr/include/sys/link.h` (Solaris), `/usr/include/elf.h` (Linux), in the *Linux Standard Base specifications* [LSB⁹], and in Drepper's paper [SYM-VER¹⁰].

A binary file intended to be specific to Linux shall set the `EI_OSABI` field to the value required by Linux [LSB¹¹].

4.1.2 Program Linkage Table (PLT) Sequences and Usage Models

Symbols for which a PLT entry must be generated

A PLT entry implements a long-branch to a destination outside of this executable file. In general, the static linker knows only the name of the destination. It does not know its address. Such a location is called an imported location or imported symbol.

SysV-based Dynamic Shared Objects (DSOs) (e.g. for Linux) also require functions exported from an executable file to have PLT entries. In effect, exported functions are treated as if they were imported, so that their definitions can be overridden (pre-empted) at dynamic link time.

A linker must generate a PLT entry for each candidate symbol cited by a relocation directive that relocates an AArch64 B/BL-class instruction (*Call and Jump relocations* (page 25)). For a Linux/SysV DSO, each STB_GLOBAL symbol with STV_DEFAULT visibility is a candidate.

Overview of PLT entry code generation

A PLT entry must be able to branch any distance. This is typically achieved by loading the destination address from the corresponding Global Object Table (GOT) entry.

⁷ <http://www.linuxbase.org/>

⁸ <http://www.akkadia.org/drepper/symbol-versioning>

⁹ <http://www.linuxbase.org/>

¹⁰ <http://www.akkadia.org/drepper/symbol-versioning>

¹¹ <http://www.linuxbase.org/>

On-demand dynamic linking constrains the code sequences that can be generated for a PLT entry. Specifically, there is a requirement from the dynamic linker for certain registers to contain certain values. Typically these are:

- The address or index of the of not-yet-linked PLT entry.
- The return address of the call to the PLT entry.

The register interface to the dynamic linker is specified by the host operating system.

OBJECT FILES

5.1 Introduction

5.1.1 Registered Vendor Names

Various symbols and names may require a vendor-specific name to avoid the potential for name-space conflicts. The list of currently registered vendors and their preferred short-hand name is given in Table 4-1, Registered Vendors. Tools developers not listed are requested to co-ordinate with Arm to avoid the potential for conflicts.

Table 5.1: Table 4-1, Registered Vendors

Name	Vendor
aeabi	Reserved to the ABI for the Arm Architecture (EABI pseudo-vendor)
AnonXyz anonXyz	Reserved to private experiments by the Xyz vendor. Guaranteed not to clash with any registered vendor name.
ARM	Arm Ltd (Note: the company, not the processor).
cx	C++ ABI pseudo-vendor
FSL	Freescall Semiconductor Inc.
GHS	Green Hills Systems
gnu	GNU compilers and tools (Free Software Foundation)
iar	IAR Systems
intel	Intel Corporation
ixs	Intel Xscale
llvm	The LLVM/Clang projects
PSI	PalmSource Inc.
RAL	Rowley Associates Ltd
somn	SOMNIUM Technologies Limited.
TASKING	Altium Ltd.
TI	TI Inc.
tls	Reserved for use in thread-local storage routines.
WRS	Wind River Systems.

To register a vendor prefix with Arm, please E-mail your request to arm.eabi at arm.com.

5.2 ELF Header

The ELF header provides a number of fields that assist in interpretation of the file. Most of these are specified in the base standard. The following fields have Arm-specific meanings.

e_machine An object file conforming to this specification must have the value EM_AARCH64 (183, 0xB7).

e_entry The base ELF specification requires this field to be zero if an application does not have an entry point. Nonetheless, some applications may require an entry point of zero (for example, via a reset vector).

A platform standard may specify that an executable file always has an entry point, in which case `e_entry` specifies that entry point, even if zero.

e_flags There are no processor-specific flags so this field shall contain zero.

5.2.1 ELF Identification

The 16-byte ELF identification (`e_ident`) provides information on how to interpret the file itself. The following values shall be used on Arm systems

EI_CLASS For object files (executable, shared and relocatable) the **EI_CLASS** shall be:

- **ELFCLASS64** for an ELF64 object file.
- **ELFCLASS32** for an ELF32 object file (**Beta**).

EI_DATA This field may be either **ELFDATA2LSB** or **ELFDATA2MSB**. The choice will be governed by the default data order in the execution environment.

I_OSABI This field shall be zero unless the file uses objects that have flags which have OS-specific meanings (for example, it makes use of a section index in the range **SHN_LOOS** through **SHN_HIOS**).

5.3 Sections

5.3.1 Special Section Indexes

No processor-specific special section indexes are defined. All processor-specific values are reserved to future revisions of this specification.

5.3.2 Section Types

The defined processor-specific section types are listed in Table 4-2, Processor specific section types. All other processor-specific values are reserved to future revisions of this specification.

Table 5.2: Table 4-2, Processor specific section types

Name	Value	Comment
SHT_AARCH64_ATTRIBUTES	0x70000003	Reserved for Object file compatibility attributes

5.3.3 Section Attribute Flags

There are no processor-specific section attribute flags defined. All processor-specific values are reserved to future revisions of this specification.

Merging of objects in sections with **SHF_MERGE**

In a section with the **SHF_MERGE** flag set, duplicate used objects may be merged and unused objects may be removed. An object is used if:

- A relocation directive addresses the object via the section symbol with a suitable addend to point to the object.
- A relocation directive addresses a symbol within the section. The used object is the one addressed by the symbol irrespective of the addend used.

5.3.4 Special Sections

Table 4-3, AArch64 special sections lists the special sections defined by this ABI.

Table 5.3: Table 4-3, AArch64 special sections

Name	Type	Attributes
<code>.ARM.attributes</code>	<code>SHT_AARCH64_ATTRIBUTES</code>	<code>none</code>
<code>.note.gnu.property</code>	<code>SHT_NOTE</code>	<code>SHF_ALLOC</code>

`.ARM.attributes` names a section that contains build attributes. See [Build Attributes](#) (page 14).

`.note.gnu.property` names a section that holds a program property note. See [\[LINUX_ABI¹²\]](#) for more information.

Additional special sections may be required by some platforms standards.

5.3.5 Section Alignment

There is no minimum alignment required for a section. Sections containing code must be at least 4-byte aligned. Platform standards may set a limit on the maximum alignment that they can guarantee (normally the minimum page size supported by the platform).

5.3.6 Build Attributes

Build attributes are encoded in a section of type `SHT_AARCH64_ATTRIBUTES`, and name `.ARM.attributes`.

Build attributes are unnecessary when a platform ABI operating system is fully specified. At this time no public build attributes have been defined for AArch64, however, software development tools are free to use attributes privately. For an introduction to AArch32 build attributes see [\[Addenda32\]](#).

5.4 String Table

There are no processor-specific extensions to the string table.

5.5 Symbol Table

There are no processor-specific symbol types or symbol bindings. All processor-specific values are reserved to future revisions of this specification.

5.5.1 st_other Values

The `st_other` member of a symbol table entry specifies the symbol's visibility in the lowest 2 bits. The top 6 bits are unused in the generic ELF ABI [\[SCO-ELF¹³\]](#), and while there are no values reserved for processor-specific semantics, many other architectures have used these bits.

The defined processor-specific `st_other` flag values are listed in [Table 4-5-1](#) (page 15).

¹² <https://github.com/hjl-tools/linux-abi/wiki>

¹³ <http://www.sco.com/developers/gabi/>

Table 5.4: Table 4-5-1, Processor specific st_other flags

Name	Mask	Comment
STO_AARCH64_VARIANT_PCS	0x80	The function associated with the symbol may follow a variant procedure call standard with different register usage convention.

A symbol table entry that is marked with the STO_AARCH64_VARIANT_PCS flag set in its st_other field may be associated with a function that follows a variant procedure call standard with different register usage convention from the one defined in the base procedure call standard for the list of argument, caller-saved and callee-saved registers [AAPCS64]. The rules in the [Call and Jump relocations](#) (page 25) section still apply to such functions. If a subroutine is called via a symbol reference that is marked with STO_AARCH64_VARIANT_PCS, then code that runs between the calling routine and the called subroutine must preserve the contents of all registers except for IP0, IP1, and the condition code flags [AAPCS64].

Static linkers must preserve the marking and propagate it to the dynamic symbol table if any reference or definition of the symbol is marked with STO_AARCH64_VARIANT_PCS, and add a DT_AARCH64_VARIANT_PCS dynamic tag if required by the [Dynamic Section](#) (page 35) section.

Note: In particular, when a call is made via the PLT entry of a symbol marked with STO_AARCH64_VARIANT_PCS, a dynamic linker cannot assume that the call follows the register usage convention of the base procedure call standard.

An example of a function that follows a variant procedure call standard with different register usage convention is one that takes parameters in scalable vector or predicate registers.

5.6 Weak Symbols

There are two forms of weak symbol:

- A weak reference — This is denoted by:
 - st_shndx=SHN_UNDEF, ELF64_ST_BIND()=STB_WEAK.
 - st_shndx=SHN_UNDEF, ELF32_ST_BIND()=STB_WEAK (**Beta**).
- A weak definition — This is denoted by:
 - st_shndx!=SHN_UNDEF, ELF64_ST_BIND()=STB_WEAK.
 - st_shndx!=SHN_UNDEF, ELF32_ST_BIND()=STB_WEAK (**Beta**).

5.6.1 Weak References

Libraries are not searched to resolve weak references. It is not an error for a weak reference to remain unsatisfied.

During linking, the symbol value of an undefined weak reference is:

- Zero if the relocation type is absolute
- The address of the place if the relocation type is pc-relative.

See [Relocation](#) (page 17) for further details.

5.6.2 Weak Definitions

A weak definition does not change the rules by which object files are selected from libraries. However, if a link set contains both a weak definition and a non-weak definition, the non-weak definition will always be used.

5.6.3 Symbol Types

All code symbols exported from an object file (symbols with binding STB_GLOBAL) shall have type STT_FUNC. All extern data objects shall have type STT_OBJECT. No STB_GLOBAL data symbol shall have type STT_FUNC. The type of an undefined symbol shall be STT_NOTYPE or the type of its expected definition.

The type of any other symbol defined in an executable section can be STT_NOTYPE. A linker is only required to provide long-branch and PLT support for symbols of type STT_FUNC.

5.6.4 Symbol names

A symbol that names a C or assembly language entity should have the name of that entity. For example, a C function called `calculate` generates a symbol called `calculate` (not `_calculate`).

Symbol names are case sensitive and are matched exactly by linkers.

Any symbol with binding STB_LOCAL may be removed from an object and replaced with an offset from another symbol in the same section under the following conditions:

- The original symbol and replacement symbol are not of type STT_FUNC, or both symbols are of type STT_FUNC.
- All relocations referring to the symbol can accommodate the adjustment in the addend field (it is permitted to convert a REL type relocation to a RELA type relocation).
- The symbol is not described by the debug information.
- The symbol is not a mapping symbol (*Mapping symbols* (page 17)).
- The resulting object, or image, is not required to preserve accurate symbol information to permit de-compilation or other post-linking optimization techniques.
- If the symbol labels an object in a section with the SHF_MERGE flag set, the relocation using symbol may be changed to use the section symbol only if the initial addend of the relocation is zero.

No tool is required to perform the above transformations; an object consumer must be prepared to do this itself if it might find the additional symbols confusing.

Note: Multiple conventions exist for the names of compiler temporary symbols (for example, ARMCC uses `Lxxx.yyy`, while GNU tools use `.Lxxx`).

Reserved symbol names

The following symbols are reserved to this and future revisions of this specification:

- Local symbols (STB_LOCAL) beginning with ‘\$’
- Symbols matching the pattern `non-empty-prefix$$non-empty-suffix`.
- Global symbols (STB_GLOBAL, STB_WEAK) beginning with ‘__aeabi_’ (double ‘_’ at start).

Note: Global symbols beginning with ‘__vendor_’ (double ‘_’ at start), where vendor is listed in *Registered Vendor Names* (page 12) are reserved to the named vendor for the purpose of providing vendor-specific toolchain support functions.

5.6.5 Mapping symbols

A section of an ELF file can contain a mixture of A64 code and data. There are inline transitions between code and data at literal pool boundaries.

Linkers, file decoders and other tools need to map binaries correctly. To support this, a number of symbols, termed mapping symbols appear in the symbol table to label the start of each sequence of bytes of the appropriate class. All mapping symbols have type `STT_NOTYPE` and binding `STB_LOCAL`. The `st_size` field is unused and must be zero.

The mapping symbols are defined in Table 4-4, Mapping symbols. It is an error for a relocation to reference a mapping symbol. Two forms of mapping symbol are supported:

- A short form that uses a dollar character and a single letter denoting the class. This form can be used when an object producer creates mapping symbols automatically. Its use minimizes string table size.
- A longer form in which the short form is extended with a period and then any sequence of characters that are legal for a symbol. This form can be used when assembler files have to be annotated manually and the assembler does not support multiple definitions of symbols.

Mapping symbols defined in a section (relocatable view) or segment (executable view) define a sequence of half-open intervals that cover the address range of the section or segment. Each interval starts at the address defined by the mapping symbol, and continues up to, but not including, the address defined by the next (in address order) mapping symbol or the end of the section or segment. A section that contains instructions must have a mapping symbol defined at the beginning of the section. If a section contains only data no mapping symbol is required. A platform ABI should specify whether or not mapping symbols are present in the executable view; they will never be present in a stripped executable file.

Table 5.5: Table 4-4, Mapping symbols

Name	Meaning
\$x \$x.<any...>	Start of a sequence of A64 instructions
\$d \$d.<any...>	Start of a sequence of data items (for example, a literal pool)

5.7 Relocation

Relocation information is used by linkers to bind symbols to addresses that could not be determined when the binary file was generated. Relocations are classified as *Static* or *Dynamic*.

- A *static relocation* relocates a place in an ELF relocatable file (`e_type = ET_REL`); a static linker processes it.
- A *dynamic relocation* is designed to relocate a place in an ELF executable file or dynamic shared object (`e_type = ET_EXEC, ET_DYN`) and to be handled by a dynamic linker, program loader, or other post-linking tool (dynamic linker henceforth).
- A dynamic linker need only process dynamic relocations; a static linker must handle any defined relocation.
- Dynamic relocations are designed to be processed quickly.
 - There are a small number of dynamic relocations whose codes are contiguous.

- Dynamic relocations relocate simple places and do not need complex field extraction or insertion.
- A static linker either:
 - Fully resolves a relocation directive.
 - Or, generates a dynamic relocation from it for processing by a dynamic linker.
- A well-formed executable file or dynamic shared object has no static relocations after static linking.

5.7.1 Relocation codes

The relocation codes for AArch64 are divided into four categories:

- Mandatory relocations that must be supported by all static linkers.
- Platform-specific relocations required by specific platform ABIs.
- Private relocations that are guaranteed never to be allocated in future revisions of this specification, but which must never be used in portable object files.
- Unallocated relocations that are reserved for use in future revisions of this specification.

5.7.2 Addends and PC-bias

A binary file may use REL or RELA relocations or a mixture of the two (but multiple relocations of the same place must use only one type).

The initial addend for a REL-type relocation is formed according to the following rules.

- If the relocation relocates data (*Static Data relocations* (page 20)) the initial value in the place is sign-extended to 64 bits.
- If the relocation relocates an instruction the immediate field of the instruction is extracted, scaled as required by the instruction field encoding, and sign-extended to 64 bits.

A RELA format relocation must be used if the initial addend cannot be encoded in the place.

There is no PC bias to accommodate in the relocation of a place containing an instruction that formulates a PC- relative address. The program counter reflects the address of the currently executing instruction.

5.7.3 Relocation types

Tables in the following sections list the relocation codes for AArch64 and record the following.

- The relocation code which is stored in the ELF64_R_TYPE or ELF32_R_TYPE component of the `r_info` field.
- The preferred mnemonic name for the relocation. This has no significance in a binary file.
- The relocation operation required. This field describes how a symbol and addend are processed by a linker. It does not describe how an initial addend value is extracted from a place (*Addends and PC-bias* (page 18)) or how the resulting relocated value is inserted or encoded into a place.
- A comment describing the kind of place that can be relocated, the part of the result value inserted into the place, and whether or not field overflow should be checked.

Relocation names and class

A mnemonic name class is used to distinguish between ELF64 and ELF32 relocation names.

- ELF64 relocations have <CLS> = AARCH64, e.g. R_AARCH64_ABS32

- ELF32 relocations have `<CLS> = AARCH64_P32`, where P32 denotes the pointer size, e.g. `R_AARCH64_P32_ABS32` (**Beta**)

Note: Within this document `<CLS>` is not expanded in instances where only a single relocation name exists.

Relocation codes

References to relocation codes are disambiguated in the following way:

- ELF64 relocation codes are bounded by parentheses: ().
- ELF32 relocation codes are bounded by brackets: [].

Static relocation codes for ELF64 object files begin at (257); dynamic ones at (1024). Both (0) and (256) should be accepted as values of `R_AARCH64_NONE`, the null relocation.

Static relocation codes for ELF32 object files begin at [1]; dynamic ones at [180].

All unallocated type codes are reserved for future allocation.

Relocation operations

The following nomenclature is used in the descriptions of relocation operations:

- S (when used on its own) is the address of the symbol.
- A is the addend for the relocation.
- P is the address of the place being relocated (derived from `r_offset`).
- X is the result of a relocation operation, before any masking or bit-selection operation is applied
- `Page(expr)` is the page address of the expression `expr`, defined as `(expr & ~0xFFF)`. (This applies even if the machine page size supported by the platform has a different value.)
- GOT is the address of the Global Offset Table, the table of code and data addresses to be resolved at dynamic link time. The GOT and each entry in it must be, 64-bit aligned for ELF64 or 32-bit aligned for ELF32.
- `GDAT(S+A)` represents a pointer-sized entry in the GOT for address `S+A`. The entry will be relocated at run time with relocation `R_<CLS>_GLOB_DAT(S+A)`.
- `G(expr)` is the address of the GOT entry for the expression `expr`.
- `Delta(S)` if S is a normal symbol, resolves to the difference between the static link address of S and the execution address of S. If S is the null symbol (ELF symbol index 0), resolves to the difference between the static link address of P and the execution address of P.
- `Indirect(expr)` represents the result of calling `expr` as a function. The result is the return value from the function that is returned in `r0`. The arguments passed to the function are defined by the platform ABI.
- `[msb:lsb]` is a bit-mask operation representing the selection of bits in a value. The bits selected range from `lsb` up to `msb` inclusive. For example, 'bits [3:0]' represents the bits under the mask `0x0000000F`. When range checking is applied to a value, it is applied before the masking operation is performed.

The value written into a target field is always reduced to fit the field. It is Q-o-I whether a linker generates a diagnostic when a relocated value overflows its target field.

Relocation types whose names end with “_NC” are non-checking relocation types. These must not generate diagnostics in case of field overflow. Usually, a non-checking type relocates an instruction that computes one of the less significant parts of a single value computed by a group of instructions (*Group*)

relocations (page 26)). Only the instruction computing the most significant part of the value can be checked for field overflow because, in general, a relocated value will overflow the fields of instructions computing the less significant parts. Some non-checking relocations may, however, be expected to check for correct alignment of the result; the notes explain when this is permitted. In ELF32 relocations an overflow check of $-2^{31} \leq X < 2^{31}$ or $0 \leq X < 2^{31}$ is equivalent to no check (i.e. 'None').

In ELF32 (**Beta**) relocations additional care must be taken when relocating an ADRP instruction which effectively uses a signed 33-bit PC-relative offset to generate a 32-bit address. The following relocations apply to ADRP:

```
R_<CLS>_ADR_PREL_PG_HI21,
R_<CLS>_ADR_GOT_PAGE,
R_<CLS>_TLSGD_ADR_PAGE21,
R_<CLS>_TSLD_ADR_PAGE21,
R_<CLS>_TLSIE_ADR_GOTTPREL_PAGE21,
R_<CLS>_TSDDESC_ADR_PAGE21
```

5.7.4 Static miscellaneous relocations

R_<CLS>_NONE (null relocation code) records that the section containing the place to be relocated depends on the section defining the symbol mentioned in the relocation directive in a way otherwise invisible to a static linker. The effect is to prevent removal of sections that might otherwise appear to be unused.

Table 5.6: Table 4-5, Null relocation codes

ELF64 Code	ELF32 Code	Name	Operation	Comment
0	0	R_<CLS>_NONE	None	
256	-	withdrawn	None	Treat as R_<CLS>_NONE.

5.7.5 Static Data relocations

See also Table 4-13, GOT-relative data relocations.

Table 5.7: Table 4-6, Data relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
257	-	R_<CLS>_ABS64	S + A	No overflow check
258	1	R_<CLS>_ABS32	S + A	Check that $-2^{31} \leq X < 2^{32}$
259	2	R_<CLS>_ABS16	S + A	Check that $-2^{15} \leq X < 2^{16}$
260	-	R_<CLS>_PREL64	S + A - P	No overflow check
261	3	R_<CLS>_PREL32	S + A - P	Check that $-2^{31} \leq X < 2^{32}$
262	4	R_<CLS>_PREL16	S + A - P	Check that $-2^{15} \leq X < 2^{16}$
314	29	R_<CLS>_PLT32	S + A - P	Check that $-2^{31} \leq X < 2^{31}$ see call and jump relocations

These overflow ranges permit either signed or unsigned narrow values to be created from the intermediate result viewed as a 64-bit signed integer. If the place is intended to hold a narrow signed value and

$\text{INTn_MAX} < X \leq \text{UINTn_MAX}$, no overflow will be detected but the positive result will be interpreted as a negative value.

5.7.6 Static AArch64 relocations

The following tables record single instruction relocations and relocations that allow a group or sequence of instructions to compute a single relocated value.

Table 5.8: Table 4-7, Group relocations to create a 16-, 32-, 48-, or 64-bit unsigned data value or address inline

ELF64 Code	ELF32 Code	Name	Operation	Comment
263	5	R_<CLS>_MOVW_UABS_G0	S + A	Set a MOV[KZ] immediate field to bits [15:0] of X; check that $0 \leq X < 2^{16}$
264	6	R_<CLS>_MOVW_UABS_G0_NC	S + A	Set a MOV[KZ] immediate field to bits [15:0] of X. No overflow check
265	7	R_<CLS>_MOVW_UABS_G1	S + A	Set a MOV[KZ] immediate field to bits [31:16] of X; check that $0 \leq X < 2^{32}$
266	-	R_<CLS>_MOVW_UABS_G1_NC	S + A	Set a MOV[KZ] immediate field to bits [31:16] of X. No overflow check
267	-	R_<CLS>_MOVW_UABS_G2	S + A	Set a MOV[KZ] immediate field to bits [47:32] of X; check that $0 \leq X < 2^{48}$
268	-	R_<CLS>_MOVW_UABS_G2_NC	S + A	Set a MOV[KZ] immediate field to bits [47:32] of X. No overflow check
269	-	R_<CLS>_MOVW_UABS_G3	S + A	Set a MOV[KZ] immediate field to bits [63:48] of X (no overflow check needed)

Table 5.9: Table 4-8, Group relocations to create a 16, 32, 48, or 64 bit signed data or offset value inline

ELF64 Code	ELF32 Code	Name	Operation	Comment
270	8	R_<CLS>_MOVW_SABS_G0	S + A	Set a MOV[NZ] immediate field using bits [15:0] of X (see notes below); check $-2^{16} \leq X < 2^{16}$
271	-	R_<CLS>_MOVW_SABS_G1	S + A	Set a MOV[NZ] immediate field using bits [31:16] of X (see notes below); check $-2^{32} \leq X < 2^{32}$
272	-	R_<CLS>_MOVW_SABS_G2	S + A	Set a MOV[NZ] immediate field using bits [47:32] of X (see notes below); check $-2^{48} \leq X < 2^{48}$

Note: These checking forms relocate MOVN or MOVZ.

$X \geq 0$: Set the instruction to MOVZ and its immediate field to the selected bits of X.

$X < 0$: Set the instruction to MOVN and its immediate field to NOT (selected bits of X).

Table 5.10: Table 4-9, Relocations to generate 19, 21 and 33 bit PC-relative addresses

ELF64 Code	ELF32 Code	Name	Operation	Comment
273	9	R_<CLS>_LD_PREL_LO19	S + A - P	Set a load-literal immediate value to bits [20:2] of X; check that $-2^{20} \leq X < 2^{20}$
274	10	R_<CLS>_ADR_PREL_LO21	S + A - P	Set an ADR immediate value to bits [20:0] of X; check that $-2^{20} \leq X < 2^{20}$
275	11	R_<CLS>_ADR_PREL_PG_HI21	Page(S+A)-Page(P)	Set an ADRP immediate value to bits [32:12] of the X; check that $-2^{32} \leq X < 2^{32}$
276	-	R_<CLS>_ADR_PREL_PG_HI21_NC	Page(S+A)-Page(P)	Set an ADRP immediate value to bits [32:12] of the X. No overflow check
277	12	R_<CLS>_ADD_ABS_LO12_NC	S + A	Set an ADD immediate value to bits [11:0] of X. No overflow check. Used with relocations ADR_PREL_PG_HI21 and ADR_PREL_PG_HI21_NC
278	13	R_<CLS>_LDST8_ABS_LO12_NC	S + A	Set an LD/ST immediate value to bits [11:0] of X. No overflow check. Used with relocations ADR_PREL_PG_HI21 and ADR_PREL_PG_HI21_NC
284	14	R_<CLS>_LDST16_ABS_LO12_NC	S + A	Set an LD/ST immediate value to bits [11:1] of X. No overflow check
285	15	R_<CLS>_LDST32_ABS_LO12_NC	S + A	Set the LD/ST immediate value to bits [11:2] of X. No overflow check
286	16	R_<CLS>_LDST64_ABS_LO12_NC	S + A	Set the LD/ST immediate value to bits [11:3] of X. No overflow check
299	17	R_<CLS>_LDST128_ABS_LO12_NC	S + A	Set the LD/ST immediate value to bits [11:4] of X. No overflow check

Note: Relocations (284, 285, 286 and 299) or [14, 15, 16, 17] are intended to be used with R_<CLS>_ADR_PREL_PG_HI21 (275) or [11] so they pick out the low 12 bits of the address and, in effect, scale that by the access size. The increased address range provided by scaled addressing is not supported by these relocations because the extra range is unusable in conjunction with R_<CLS>_ADR_PREL_PG_HI21.

Although overflow must not be checked, a linker should check that the value of X is aligned to a multiple of the datum size.

Table 5.11: Table 4-10, Relocations for control-flow instructions -
all offsets are a multiple of 4

ELF64 Code	ELF32 Code	Name	Operation	Comment
279	18	R_<CLS>_TSTBR14	S+A-P	Set the immediate field of a TBZ/TBNZ instruction to bits [15:2] of X; check $-2^{15} \leq X < 2^{15}$
280	19	R_<CLS>_CONDBR19	S+A-P	Set the immediate field of a conditional branch instruction to bits [20:2] of X; check $-2^{20} \leq X < 2^{20}$
282	20	R_<CLS>_JUMP26	S+A-P	Set a B immediate field to bits [27:2] of X; check that $-2^{27} \leq X < 2^{27}$
283	21	R_<CLS>_CALL26	S+A-P	Set a CALL immediate field to bits [27:2] of X; check that $-2^{27} \leq X < 2^{27}$

Table 5.12: Table 4-11, Group relocations to create a 16, 32, 48, or
64 bit PC-relative offset inline

ELF64 Code	ELF32 Code	Name	Operation	Comment
287	22	R_<CLS>_MOVW_PREL_G0	S+A-P	Set a MOV[NZ]immediate field to bits [15:0] of X (see notes below)
288	23	R_<CLS>_MOVW_PREL_G0_NC	S+A-P	Set a MOVK immediate field to bits [15:0] of X. No overflow check
289	24	R_<CLS>_MOVW_PREL_G1	S+A-P	Set a MOV[NZ]immediate field to bits [31:16] of X (see notes below)
290	-	R_<CLS>_MOVW_PREL_G1_NC	S+A-P	Set a MOVK immediate field to bits [31:16] of X. No overflow check
291	-	R_<CLS>_MOVW_PREL_G2	S+A-P	Set a MOV[NZ]immediate value to bits [47:32] of X (see notes below)
292	-	R_<CLS>_MOVW_PREL_G2_NC	S+A-P	Set a MOVK immediate field to bits [47:32] of X. No overflow check
293	-	R_<CLS>_MOVW_PREL_G3	S+A-P	Set a MOV[NZ]immediate value to bits [63:48] of X (see notes below)

Note: Non-checking (_NC) forms relocate MOVK; checking forms relocate MOVN or MOVZ.

$X \geq 0$: Set the instruction to MOVZ and its immediate value to the selected bits of X; for relocation R_..._Gn, check in ELF64 that $X < \{G0: 2^{16}, G1: 2^{32}, G2: 2^{48}\}$ (no check for R_..._G3); in ELF32 only check $X < 2^{16}$ for R_..._G0.

$X < 0$: Set the instruction to MOVN and its immediate value to NOT (selected bits of X); for relocation R_..._Gn, check in ELF64 that $-\{G0: 2^{16}, G1: 2^{32}, G2: 2^{48}\} \leq X$ (no check for R_..._G3); in ELF32 only check that $-2^{16} \leq X$ for R_..._G0.

Table 5.13: Table 4-12, Group relocations to create a 16, 32, 48, or 64 bit GOT-relative offsets inline

ELF64 Code	ELF32 Code	Name	Operation	Comment
300	-	R_<CLS>_MOVW_GOTOFF_G0	G(GDAT(S+A)) -GOT	Set a MOV[NZ] immediate field to bits [15:0] of X (see notes below)
301	-	R_<CLS>_MOVW_GOTOFF_G0_NC	G(GDAT(S+A)) -GOT	Set a MOVK immediate field to bits [15:0] of X. No overflow check
302	-	R_<CLS>_MOVW_GOTOFF_G1	G(GDAT(S+A)) -GOT	Set a MOV[NZ] immediate value to bits [31:16] of X (see notes below)
303	-	R_<CLS>_MOVW_GOTOFF_G1_NC	G(GDAT(S+A)) -GOT	Set a MOVK immediate value to bits [31:16] of X. No overflow check
304	-	R_<CLS>_MOVW_GOTOFF_G2	G(GDAT(S+A)) -GOT	Set a MOV[NZ] immediate value to bits [47:32] of X (see notes below)
305	-	R_<CLS>_MOVW_GOTOFF_G2_NC	G(GDAT(S+A)) -GOT	Set a MOVK immediate value to bits [47:32] of X. No overflow check
306	-	R_<CLS>_MOVW_GOTOFF_G3	G(GDAT(S+A)) -GOT	Set a MOV[NZ] immediate value to bits [63:48] of X (see notes below)

Note: Non-checking (_NC) forms relocate MOVK; checking forms relocate MOVN or MOVZ.

Table 5.14: Table 4-13, GOT-relative data relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
307	-	R_<CLS>_GOTREL64	S+A-GOT	Set the data to a 64-bit offset relative to the GOT.
308	-	R_<CLS>_GOTREL32	S+A-GOT	Set the data to a 32-bit offset relative to GOT, treated as signed; check that $-2^{31} \leq X < 2^{31}$

Table 5.15: Table 4-14, GOT-relative instruction relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
309	25	R_<CLS>_GOT_LD_PREL19	G(GDAT(S+A))- P	Set a load-literal immediate field to bits [20:2] of X; check $-2^{20} \leq X < 2^{20}$
310	-	R_<CLS>_LD64_GOTOFF_LO15	G(GDAT(S+A))- GOT	Set a LD/ST immediate field to bits [14:3] of X; check that $0 \leq X < 2^{15}$, $X \& 7 = 0$
311	26	R_<CLS>_ADR_GOT_PAGE	Page(G(GDAT(S+A)))- Page(P)	Set the immediate value of an ADRP to bits [32:12] of X; check that $-2^{32} \leq X < 2^{32}$

Continued on next page

Table 5.15 – continued from previous page

ELF64 Code	ELF32 Code	Name	Operation	Comment
312	-	R_<CLS>_LD64_GOT_LO12_NC	G(GDAT(S+A))	Set the LD/ST immediate field to bits [11:3] of X. No overflow check; check that $X \& 7 = 0$
-	27	R_<CLS>_LD32_GOT_LO12_NC	G(GDAT(S+A))	Set the LD/ST immediate field to bits [11:2] of X. No overflow check; check that $X \& 3 = 0$
313	-	R_<CLS>_LD64_GOTPAGE_LO15	G(GDAT(S+A))-Page(GOT)	Set the LD/ST immediate field to bits [14:3] of X; check that $0 \leq X < 2^{15}$, $X \& 7 = 0$
-	28	R_<CLS>_LD32_GOTPAGE_LO14	G(GDAT(S+A))-Page(GOT)	Set the LD/ST immediate field to bits [13:2] of X; check that $0 \leq X < 2^{14}$, $X \& 3 = 0$

5.7.7 Call and Jump relocations

There is one relocation code (R_<CLS>_CALL26) for function call (BL) instructions and one (R_<CLS>_JUMP26) for jump (B) instructions. The (R_<CLS>_PLT32) relocation is a data relocation for calculating the offset to a function. This can be used as the target of an indirect jump.

A linker may use a veneer (a sequence of instructions) to implement a relocated branch if the relocation is either

R_<CLS>_CALL26, R_<CLS>_JUMP26 or R_<CLS>_PLT32 and:

- The target symbol has type STT_FUNC.
- Or, the target symbol and relocated place are in separate sections input to the linker.
- Or, the target symbol is undefined (external to the link unit).

In all other cases a linker shall diagnose an error if relocation cannot be effected without a veneer. A linker generated veneer may corrupt registers IP0 and IP1 [AAPCS64] and the condition flags, but must preserve all other registers. Linker veneers may be needed for a number of reasons, including, but not limited to:

- Target is outside the addressable span of the branch instruction (+/- 128MB).
- Target address will not be known until run time, or the target address might be pre-empted.

In some systems indirect calls may also use veneers in order to support dynamic linkage that preserves pointer comparability (all reference to the function resolve to the same address).

On platforms that do not support dynamic pre-emption of symbols, an unresolved weak reference to a symbol relocated by R_<CLS>_CALL26 shall be treated as a jump to the next instruction (the call becomes a no-op). The behaviour of R_<CLS>_JUMP26 and R_<CLS>_PLT32 in these conditions is not specified by this standard.

5.7.8 Group relocations

A relocation code whose name ends in `_Gn` or `_Gn_NC` ($n = 0, 1, 2, 3$) relocates an instruction in a group of instructions that generate a single value or address (see Table 4-7, Table 4-8, Table 4-11, Table 4-12). Each such relocation relocates one instruction in isolation, with no need to determine all members of the group at link time.

These relocations operate by performing the relocation calculation then extracting a field from the result X . Generating the field for a G_n relocation directive starts by examining the residual value Y_n after the bits of $\text{abs}(X)$ corresponding to less significant fields have been masked off from X . If M is the mask specified in the table recording the relocation directive, $Y_n = \text{abs}(X) \& \sim((M \& -M) - 1)$.

Overflow checking is performed on Y_n unless the name of the relocation ends in “_NC”.

Finally the bit-field of X specified in the table (those bits of X picked out by 1-bits in M) is encoded into the instruction’s literal field as specified in the table. In some cases other instruction bits may need to be changed according to the sign of X .

For “MOVW” type relocations it is the assembler’s responsibility to encode the hw bits (bits 21 and 22) to indicate the bits in the target value that the immediate field represents.

5.7.9 Proxy-generating relocations

A number of relocations generate proxy locations that are then subject to dynamic relocation. The proxies are normally gathered together in a single table, called the Global Offset Table or GOT. Table 4-12, Group relocations to create a 16, 32, 48, or 64 bit GOT-relative offsets inline and Table 4-14, GOT-relative instruction relocations list the relocations that generate proxy entries.

All of the GOT entries generated by these relocations are subject to dynamic relocations (*Dynamic relocations* (page 31)).

5.7.10 Relocations for thread-local storage

The static relocations needed to support thread-local storage in a SysV-type environment are listed in tables in the following subsections

In addition to the terms defined in *Relocation types* (page 18), the tables listing the static relocations relating to thread-local storage use the following terms in the column named Operation.

- `GLDM(S)` represents a consecutive pair of pointer-sized entries in the GOT for the load module index of the symbol S . The first pointer-sized entry will be relocated with `R_<CLS>_TLS_DTPMOD(S)`; the second pointer-sized entry will contain the constant 0.
- `GTLSIDX(S,A)` represents a consecutive pair of pointer-sized entries in the GOT. The entry contains a `tls_index` structure describing the thread-local variable located at offset A from thread-local symbol S . The first pointer-sized entry will be relocated with `R_<CLS>_TLS_DTPMOD(S)`, the second pointer-sized entry will be relocated with `R_<CLS>_TLS_DTPREL(S+A)`.
- `GTPREL(S+A)` represents a pointer-sized entry in the GOT for the offset from the current thread pointer (TP) of the thread-local variable located at offset A from the symbol S . The entry will be relocated with `R_<CLS>_TLS_TPREL(S+A)`.
- `GTLSDESC(S+A)` represents a consecutive pair of pointer-sized entries in the GOT which contain a `tlsdesc` structure describing the thread-local variable located at offset A from thread-local symbol S . The first entry holds a pointer to the variable’s TLS descriptor resolver function and the second entry holds a platform-specific offset or pointer. The pair of pointer-sized entries will be relocated with `R_<CLS>_TLSDESC(S+A)`.
- `LDM(S)` resolves to the load module index of the symbol S .
- `DTPREL(S+A)` resolves to the offset from its module’s TLS block of the thread local variable located at offset A from thread-local symbol S .

- $\text{TPREL}(S+A)$ resolves to the offset from the current thread pointer (TP) of the thread local variable located at offset A from thread-local symbol S.
- $\text{TLSDESC}(S+A)$ resolves to a contiguous pair of pointer-sized values, as created by $\text{GTLSDESC}(S+A)$.

General Dynamic thread-local storage model

Note: Non-checking ($_NC$) MOVW forms relocate MOVK; checking forms relocate MOVN or MOVZ.

$X \geq 0$: Set the instruction to MOVZ and its immediate value to the selected bits of X; check that $X < 2^{32}$.

$X < 0$: Set the instruction to MOVN and its immediate value to NOT (selected bits of X); check that $-2^{32} \leq X$.

Local Dynamic thread-local storage model

Table 5.16: Table 4-16, Local Dynamic TLS relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
517	83	$R_{<CLS>_TLSLD_ADR_PREL21}$	$G(\text{GLDM}(S)) - P$	Set an ADR immediate field to bits [20:0] of X; check $-2^{20} \leq X < 2^{20}$
518	84	$R_{<CLS>_TLSLD_ADR_PAGE21}$	$\text{Page}(G(\text{GLDM}(S))) - \text{Page}(P)$	Set an ADRP immediate field to bits [32:12] of X; check $-2^{32} \leq X < 2^{32}$
519	85	$R_{<CLS>_TLSLD_ADD_LO12_NC}$	$G(\text{GLDM}(S))$	Set an ADD immediate field to bits [11:0] of X. No overflow check
520	-	$R_{<CLS>_TLSLD_MOVW_G1}$	$G(\text{GLDM}(S)) - \text{GOT}$	Set a MOV[NZ] immediate field to bits [31:16] of X (see notes below)
521	-	$R_{<CLS>_TLSLD_MOVW_G0_NC}$	$G(\text{GLDM}(S)) - \text{GOT}$	Set a MOVK immediate field to bits [15:0] of X. No overflow check
522	86	$R_{<CLS>_TLSLD_LD_PREL19}$	$G(\text{GLDM}(S)) - P$	Set a load-literal immediate field to bits [20:2] of X; check $-2^{20} \leq X < 2^{20}$
523	-	$R_{<CLS>_TLSLD_MOVW_DTPREL_G2}$	$\text{DTPREL}(S+A)$	Set a MOV[NZ] immediate field to bits [47:32] of X (see notes below)
524	87	$R_{<CLS>_TLSLD_MOVW_DTPREL_G1}$	$\text{DTPREL}(S+A)$	Set a MOV[NZ] immediate field to bits [31:16] of X (see notes below)
525	-	$R_{<CLS>_TLSLD_MOVW_DTPREL_G1_NC}$	$\text{DTPREL}(S+A)$	Set a MOVK immediate field to bits [31:16] of X. No overflow check
526	88	$R_{<CLS>_TLSLD_MOVW_DTPREL_G0}$	$\text{DTPREL}(S+A)$	Set a MOV[NZ] immediate field to bits [15:0] of X (see notes below)
527	89	$R_{<CLS>_TLSLD_MOVW_DTPREL_G0_NC}$	$\text{DTPREL}(S+A)$	Set a MOVK immediate field to bits [15:0] of X. No overflow check

Continued on next page

Table 5.16 – continued from previous page

ELF64 Code	ELF32 Code	Name	Operation	Comment
528	90	R_<CLS>_TLSLD_ADD_DTPREL_HI12	DTPREL(S+A)	Set an ADD immediate field to bits [23:12] of X; check $0 \leq X < 2^{24}$
529	91	R_<CLS>_TLSLD_ADD_DTPREL_LO12	DTPREL(S+A)	Set an ADD immediate field to bits [11:0] of X; check $0 \leq X < 2^{12}$
530	92	R_<CLS>_TLSLD_ADD_DTPREL_LO12_NC	DTPREL(S+A)	Set an ADD immediate field to bits [11:0] of X. No overflow check
531	93	R_<CLS>_TLSLD_LDST8_DTPREL_LO12	DTPREL(S+A)	Set a LD/ST offset field to bits [11:0] of X; check $0 \leq X < 2^{12}$
532	94	R_<CLS>_TLSLD_LDST8_DTPREL_LO12_NC	DTPREL(S+A)	Set a LD/ST offset field to bits [11:0] of X. No overflow check
533	95	R_<CLS>_TLSLD_LDST16_DTPREL_LO12	DTPREL(S+A)	Set a LD/ST offset field to bits [11:1] of X; check $0 \leq X < 2^{12}$
534	96	R_<CLS>_TLSLD_LDST16_DTPREL_LO12_NC	DTPREL(S+A)	Set a LD/ST offset field to bits [11:1] of X. No overflow check
535	97	R_<CLS>_TLSLD_LDST32_DTPREL_LO12	DTPREL(S+A)	Set a LD/ST offset field to bits [11:2] of X; check $0 \leq X < 2^{12}$
536	98	R_<CLS>_TLSLD_LDST32_DTPREL_LO12_NC	DTPREL(S+A)	Set a LD/ST offset field to bits [11:2] of X. No overflow check
537	99	R_<CLS>_TLSLD_LDST64_DTPREL_LO12	DTPREL(S+A)	Set a LD/ST offset field to bits [11:3] of X; check $0 \leq X < 2^{12}$
538	100	R_<CLS>_TLSLD_LDST64_DTPREL_LO12_NC	DTPREL(S+A)	Set a LD/ST offset field to bits [11:3] of X. No overflow check
572	101	R_<CLS>_TLSLD_LDST128_DTPREL_LO12	DTPREL(S+A)	Set a LD/ST offset field to bits [11:4] of X; check $0 \leq X < 2^{12}$
573	102	R_<CLS>_TLSLD_LDST128_DTPREL_LO12_NC	DTPREL(S+A)	Set a LD/ST offset field to bits [11:4] of X. No overflow check

Note: Non-checking (_NC) MOVW forms relocate MOVK; checking forms relocate MOVN or MOVZ.

$X \geq 0$: Set the instruction to MOVZ and its immediate value to the selected bits S; for relocation $R_{\dots Gn}$, check in ELF64 that $X < \{G0: 2^{16}, G1: 2^{32}, G2: 2^{48}\}$ (no check for $R_{\dots G3}$); in ELF32 only check that $X < 2^{16}$ for $R_{\dots G0}$.

$X < 0$: Set the instruction to MOVN and its immediate value to NOT (selected bits of); for relocation $R_{\dots Gn}$, check in ELF64 that $-\{G0: 2^{16}, G1: 2^{32}, G2: 2^{48}\} \leq X$ (no check for $R_{\dots G3}$); in ELF32 only check that $-2^{16} \leq X$ for $R_{\dots G0}$.

For scaled-addressing relocations (533-538, 572 and 573) or [95-102] a linker should check that X is a multiple of the datum size.

Initial Exec thread-local storage model

Table 5.17: Table 4-17, Initial Exec TLS relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
539	-	R_<CLS>_TLSIE_MOVW_GOTTPREL_G1	G(GTPREL(S+A)) – GOT	Set a MOV[NZ] immediate field to bits [31:16] of X (see notes below)
540	-	R_<CLS>_TLSIE_MOVW_GOTTPREL_G0_NC	G(GTPREL(S+A)) – GOT	Set MOVK immediate to bits [15:0] of X. No overflow check
541	103	R_<CLS>_TLSIE_ADR_GOTTPREL_PAGE21	Page(G(GTPREL(S+A))) – Page(P)	Set an ADRP immediate field to bits [32:12] of X; check $-2^{32} \leq X < 2^{32}$
542	-	R_<CLS>_TLSIE_LD64_GOTTPREL_LO12_NC	G(GTPREL(S+A))	Set an LD offset field to bits [11:3] of X. No overflow check; check that $X \& 7 = 0$
-	104	R_<CLS>_TLSIE_LD32_GOTTPREL_LO12_NC	G(GTPREL(S+A))	Set an LD offset field to bits [11:2] of X. No overflow check; check that $X \& 3 = 0$
543	105	R_<CLS>_TLSIE_LD_GOTTPREL_PREL19	G(GTPREL(S+A)) – P	Set a load-literal immediate to bits [20:2] of X; check $-2^{20} \leq X < 2^{20}$

Note: Non-checking (_NC) MOVW forms relocate MOVK; checking forms relocate MOVN or MOVZ.

Local Exec thread-local storage model

Table 5.18: Table 4-18, Local Exec TLS relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
544	-	R_<CLS>_TLSLE_MOVW_TPREL_G2	TPREL(S+A)	Set a MOV[NZ] immediate field to bits [47:32] of X (see notes below)
545	106	R_<CLS>_TLSLE_MOVW_TPREL_G1	TPREL(S+A)	Set a MOV[NZ] immediate field to bits [31:16] of X (see notes below)
546	-	R_<CLS>_TLSLE_MOVW_TPREL_G1_NC	TPREL(S+A)	Set a MOVK immediate field to bits [31:16] of X. No overflow check
547	107	R_<CLS>_TLSLE_MOVW_TPREL_G0	TPREL(S+A)	Set a MOV[NZ] immediate field to bits [15:0] of X (see notes below)
548	108	R_<CLS>_TLSLE_MOVW_TPREL_G0_NC	TPREL(S+A)	Set a MOVK immediate field to bits [15:0] of X. No overflow check
549	109	R_<CLS>_TLSLE_ADD_TPREL_HI12	TPREL(S+A)	Set an ADD immediate field to bits [23:12] of X; check $0 \leq X < 2^{24}$.
550	110	R_<CLS>_TLSLE_ADD_TPREL_LO12	TPREL(S+A)	Set an ADD immediate field to bits [11:0] of X; check $0 \leq X < 2^{12}$.

Continued on next page

Table 5.18 – continued from previous page

ELF64 Code	ELF32 Code	Name	Operation	Comment
551	111	R_<CLS>_TLSLE_ADD_TPREL_LO12_NC	TPREL(S+A)	Set an ADD immediate field to bits [11:0] of X. No overflow check
552	112	R_<CLS>_TLSLE_LDST8_TPREL_LO12	TPREL(S+A)	Set a LD/ST offset field to bits [11:0] of X; check $0 \leq X < 2^{12}$.
553	113	R_<CLS>_TLSLE_LDST8_TPREL_LO12_NC	TPREL(S+A)	Set a LD/ST offset field to bits [11:0] of X. No overflow check
554	114	R_<CLS>_TLSLE_LDST16_TPREL_LO12	TPREL(S+A)	Set a LD/ST offset field to bits [11:1] of X; check $0 \leq X < 2^{12}$
555	115	R_<CLS>_TLSLE_LDST16_TPREL_LO12_NC	TPREL(S+A)	Set a LD/ST offset field to bits [11:1] of X. No overflow check
556	116	R_<CLS>_TLSLE_LDST32_TPREL_LO12	TPREL(S+A)	Set a LD/ST offset field to bits [11:2] of X; check $0 \leq X < 2^{12}$
557	117	R_<CLS>_TLSLE_LDST32_TPREL_LO12_NC	TPREL(S+A)	Set a LD/ST offset field to bits [11:2] of X. No overflow check
558	118	R_<CLS>_TLSLE_LDST64_TPREL_LO12	TPREL(S+A)	Set a LD/ST offset field to bits [11:3] of X; check $0 \leq X < 2^{12}$
559	119	R_<CLS>_TLSLE_LDST64_TPREL_LO12_NC	TPREL(S+A)	Set a LD/ST offset field to bits [11:3] of X. No overflow check
570	120	R_<CLS>_TLSLE_LDST128_TPREL_LO12	TPREL(S+A)	Set a LD/ST offset field to bits [11:4] of X; check $0 \leq X < 2^{12}$
571	121	R_<CLS>_TLSLE_LDST128_TPREL_LO12_NC	TPREL(S+A)	Set a LD/ST offset field to bits [11:4] of X. No overflow check

Note: Non-checking (_NC) MOVW forms relocate MOVK; checking forms relocate MOVN or MOVZ.

For scaled-addressing relocations (554-559, 570 and 571) or [112-121] a linker should check that X is a multiple of the datum size.

Thread-local storage descriptors

Table 5.19: Table 4-19, TLS descriptor relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
560	122	R_<CLS>_TLSDESC_LD_PREL19	G(GTLSDESC(S+A)) – P	Set a load-literal immediate to bits [20:2]; check $-220 \leq X < 2^{20}$; check $X \& 3 = 0$.
561	123	R_<CLS>_TLSDESC_ADR_PREL21	G(GTLSDESC(S+A)) – P	Set an ADR immediate field to bits [20:0]; check $-220 \leq X < 2^{20}$.
562	124	R_<CLS>_TLSDESC_ADR_PAGE21	Page(G(GTLSDESC(S+A))) – Page(P)	Set an ADRP immediate field to bits [32:12] of X; check $-2^{32} \leq X < 2^{32}$.
563	-	R_<CLS>_TLSDESC_LD64_LO12	G(GTLSDESC(S+A))	Set an LD offset field to bits [11:3] of X. No overflow check; check $X \& 7 = 0$.
-	125	R_<CLS>_TLSDESC_LD32_LO12	G(GTLSDESC(S+A))	Set an LD offset field to bits [11:2] of X. No overflow check; check $X \& 3 = 0$.

Continued on next page

Table 5.19 – continued from previous page

ELF64 Code	ELF32 Code	Name	Operation	Comment
564	126	R_<CLS>_TLSDESC_ADD_LO12	G(GTLSDESC(S+A))	Set an ADD immediate field to bits [11:0] of X. No overflow check.
565	-	R_<CLS>_TLSDESC_OFF_G1	G(GTLSDESC(S+A)) – GOT	Set a MOV[NZ] immediate field to bits [31:16] of X; check $-2^{32} \leq X < 2^{32}$. See notes below.
566	-	R_<CLS>_TLSDESC_OFF_G0_NC	G(GTLSDESC(S+A)) – GOT	Set a MOVK immediate field to bits [15:0] of X. No overflow check.
567	-	R_<CLS>_TLSDESC_LDR	None	For relaxation only. Must be used to identify an LDR instruction which loads the TLS descriptor function pointer for S + A if it has no other relocation.
568	-	R_<CLS>_TLSDESC_ADD	None	For relaxation only. Must be used to identify an ADD instruction which computes the address of the TLS Descriptor for S + A if it has no other relocation.
569	127	R_<CLS>_TLSDESC_CALL	None	For relaxation only. Must be used to identify a BLR instruction which performs an indirect call to the TLS descriptor function for S + A.

Note: $X \geq 0$: Set the instruction to MOVZ and its immediate value to the selected bits of X.

$X < 0$: Set the instruction to MOVN and its immediate value to NOT (selected bits of X).

Relocation codes R_<CLS>_TLSDESC_LDR, R_<CLS>_TLSDESC_ADD and R_<CLS>_TLSDESC_CALL are needed to permit linker optimization of TLS descriptor code sequences to use Initial-exec or Local-exec TLS sequences; this can only be done if all relevant uses of TLS descriptors are marked to permit accurate relaxation. Object producers that are unable to satisfy this requirement must generate traditional General-dynamic TLS sequences using the relocations described in *General Dynamic thread-local storage model* (page 27). The details of TLS descriptors are beyond the scope of this specification; a general introduction can be found in [TLSDESC¹⁴].

5.7.11 Dynamic relocations

The dynamic relocations for those execution environments that support only a limited number of run-time relocation types are listed in Table 4-20, Dynamic relocations. The enumeration of dynamic relocations commences at (1024) or [180] and the range is compact.

¹⁴ <http://www.fsfla.org/~lxoliva/writeups/TLS/paper-lk2006.pdf>

Table 5.20: Table 4-20, Dynamic relocations

ELF64 Code	ELF32 Code	Name	Operation	Comment
1024	180	R_<CLS>_COPY		See note below.
1025	181	R_<CLS>_GLOB_DAT	S + A	See note below
1026	182	R_<CLS>_JUMP_SLOT	S + A	See note below
1027	183	R_<CLS>_RELATIVE	Delta(S) + A	See note below
1028	184	R_<CLS>_TLS_IMPDEF1		See note below
1029	185	R_<CLS>_TLS_IMPDEF2		See note below
		R_<CLS>_TLS_DTPREL	DTPREL(S+A)	See note below
		R_<CLS>_TLS_DTPMOD	LDM(S)	See note below
1030	186	R_<CLS>_TLS_TPREL	TPREL(S+A)	
1031	187	R_<CLS>_TLSDESC	TLSDESC(S+A)	Identifies a TLS descriptor to be filled
1032	188	R_<CLS>_IRELATIVE	Indirect(Delta(S) + A)	See note below.

With the exception of R_<CLS>_COPY all dynamic relocations require that the place being relocated is an 8-byte aligned 64-bit data location in ELF64 or a 4-byte aligned 32-bit data location in ELF32.

R_<CLS>_COPY may only appear in executable ELF files where e_type is set to ET_EXEC. The effect is to cause the dynamic linker to locate the target symbol in a shared library object and then to copy the number of bytes specified by its st_size field to the place. The address of the place is then used to pre-empt all other references to the specified symbol. It is an error if the storage space allocated in the executable is insufficient to hold the full copy of the symbol. If the object being copied contains dynamic relocations then the effect must be as if those relocations were performed before the copy was made.

R_<CLS>_COPY is normally only used in SysV type environments where the executable is not position-independent and references by the code and read-only data sections cannot be relocated dynamically to refer to an object that is defined in a shared library.

The need for copy relocations can be avoided if a compiler generates all code references to such objects indirectly through a dynamically relocatable location and if all static data references are placed in relocatable regions of the image. In practice, this is difficult to achieve without source-code annotation. A better approach is to avoid defining static global data in shared libraries.

R_<CLS>_GLOB_DAT relocates a GOT entry used to hold the address of a (data) symbol which must be resolved at load time.

R_<CLS>_JUMP_SLOT is used to mark code targets that will be executed.

- On platforms that support dynamic binding the relocations may be performed lazily on demand.
- The initial value stored in the place is the offset to the entry sequence stub for the dynamic linker. It must be adjusted during initial loading by the offset of the load address of the segment from its link address.
- Addresses stored in the place of these relocations may not be used for pointer comparison until the relocation after has been resolved.
- Because the initial value of the place is not related to the ultimate target of a R_<CLS>_JUMP_SLOT relocation the addend A of such a REL-type relocation shall be zero rather than the initial content of the place. A platform ABI shall prescribe whether or not the r_addend field of such a RELA-type relocation is honored. (There may be security-related reasons not to do so).

R_<CLS>_RELATIVE represents a relative adjustment to the place based on the load address of the object relative to its original link address. All symbols defined in the same segment will have the same relative adjustment. If S is the null symbol (ELF symbol index 0) then the adjustment is based on the segment defining the place. On systems where all segments are mapped contiguously the adjustment will be the same for each relocation, thus adjustment never needs to resolve the symbol. This relocation represents an optimization; it can be used to replace R_<CLS>_GLOB_DAT when the symbol resolves to the current dynamic shared object.

R_<CLS>_IRELATIVE represents a dynamic selection of the place's resolved value. The means by which this relocation is generated is platform specific, as are the conditions that must hold when resolving takes place.

Relocations R_AARCH64_TLS_DTPREL, R_AARCH64_TLS_DTPMOD and R_AARCH64_TLS_TPREL were previously documented as R_AARCH64_TLS_DTPREL64, R_AARCH64_TLS_DTPMOD64 and R_AARCH64_TLS_TPREL64 respectively. The old names can be supported if needed for backwards compatibility.

It is implementation defined whether R_<CLS>_TLS_IMPDEF1 implements R_<CLS>_TLS_DTPREL and R_<CLS>_TLS_IMPDEF2 implements R_<CLS>_TLS_DTPMOD or whether R_<CLS>_TLS_IMPDEF1 implements R_<CLS>_TLS_DTPMOD and R_<CLS>_TLS_IMPDEF2 implements R_<CLS>_TLS_DTPREL; a platform must document its choice¹⁷.

5.7.12 Private and platform-specific relocations

Private relocations for vendor experiments:

- 0xE000 to 0xEFFF for ELF64
- 0xE0 to 0xEF for ELF32

Platform ABI defined relocations:

- 0xF000 to 0xFFFF for ELF64
- 0xF0 to 0xFF for ELF32

Platform ABI relocations can only be interpreted when the EI_OSABI field is set to indicate the Platform ABI governing the definition.

All of the above codes will not be assigned by any future version of this standard.

5.7.13 Unallocated relocations

All unallocated relocation types are reserved for use by future revisions of this specification.

5.7.14 Idempotency

All RELA type relocations are idempotent. They may be reapplied to the place and the result will be the same. This allows a static linker to preserve full relocation information for an image by converting all REL type relocations into RELA type relocations.

Note: A REL type relocation can only be idempotent if the original addend was zero and if subsequent re-linking assumes that REL relocations have zero for all addends.

¹⁷ Earlier versions of this specification required that R_<CLS>_TLS_IMPDEF1 implement R_<CLS>_TLS_DTPREL and R_<CLS>_TLS_IMPDEF2 implement R_<CLS>_TLS_DTPMOD; however the Linux platform ABI has always implemented the alternative specification. It is recommended that new platforms follow the Linux platform specification as this is the most widely adopted.

PROGRAM LOADING AND DYNAMIC LINKING

This section provides details of AArch64-specific definitions and changes relating to executable images.

6.1 Program Header

The Program Header provides a number of fields that assist in interpretation of the file. Most of these are specified in the base standard [SCO-ELF¹⁵]. The following fields have AArch64-specific meanings.

p_type Table 5-1, Processor-specific segment types lists the processor-specific segment types.

Table 6.1: Table 5-1, Processor-specific segment types

Name	p_type	Meaning
PT_AARCH64_ARCHEXT	0x70000000	Reserved for architecture compatibility information
PT_AARCH64_UNWIND	0x70000001	Reserved for exception unwinding tables

A segment of type PT_AARCH64_ARCHEXT (if present) contains information describing the architecture capabilities required by the executable file. Not all platform ABIs require this segment; the Linux ABI does not. If the segment is present it must appear before segment of type PT_LOAD.

PT_AARCH64_UNWIND (if present) describes the location of a program's exception unwind tables.

p_flags There are no AArch64-specific flags.

6.1.1 Platform architecture compatibility data

At this time this ABI specifies no generic platform architecture compatibility data.

6.2 Program Property

The following processor-specific program property types [LINUX_ABI¹⁶] are defined:

Table 6.2: Table 5-2, Program Property Type

Name	Value
GNU_PROPERTY_AARCH64_FEATURE_1_AND	0xc0000000

GNU_PROPERTY_AARCH64_FEATURE_1_AND describes a set of processor features with which an ELF object or executable image is compatible, but does not require in order to execute correctly. It has a single 32-bit value for the pr_data field. Each bit represents a separate feature.

¹⁵ <http://www.sco.com/developers/gabi/>

¹⁶ <https://github.com/hjl-tools/linux-abi/wiki>

Static linkers processing ELF relocatable objects must set the feature bit in the output object or image only if all the input objects have the corresponding feature bit set. For each feature bit set in an ELF executable or shared library, a loader may enable the corresponding processor feature for that ELF file.

The following bits are defined for GNU_PROPERTY_AARCH64_FEATURE_1_AND:

Table 6.3: Table 5-3, GNU_PROPERTY_AARCH64_FEATURE_1_AND Bit Flags

Name	Value
GNU_PROPERTY_AARCH64_FEATURE_1_BTI	1U << 0
GNU_PROPERTY_AARCH64_FEATURE_1_PAC	1U << 1

GNU_PROPERTY_AARCH64_FEATURE_1_BTI This indicates that all executable sections are compatible with Branch Target Identification mechanism. An executable or shared object with this bit set is required to generate *Custom PLTs* (page 35) with BTI instruction.

GNU_PROPERTY_AARCH64_FEATURE_1_PAC This indicates that all executable sections have Return Address Signing enabled. An executable or shared object with this bit set can generate *Custom PLTs* (page 35) with a PAC instruction.

6.3 Program Loading

6.3.1 Process GNU_PROPERTY_AARCH64_FEATURE_1_BTI

If Branch Target Identification mechanism is enabled on a processor then the Guard Page (GP) bit must be disabled on the memory image of loaded executable segments of executables and shared objects that do not have GNU_PROPERTY_AARCH64_FEATURE_1_BTI set, before execution is transferred to them.

6.4 Dynamic Linking

6.4.1 Custom PLTs

- To support Branch Target Identification mechanism, in the presense of a GNU_PROPERTY_AARCH64_FEATURE_1_BTI all PLT entries generated by the linker must have a BTI instruction as the first instruction. The linker must add the DT_AARCH64_BTI_PLT (*Table 5-4, AArch64 specific dynamic array tags* (page 35)) tag to the dynamic section.
- To support Pointer Authentication, PLT entries generated by the linker can have an authenticating instruction as the final instruction before branching back. The linker must add the DT_AARCH64_PAC_PLT (*Table 5-4, AArch64 specific dynamic array tags* (page 35)) tag to the dynamic section.
- If the linker generates custom PLT entries with both BTI and PAC instructions, it must add both DT_AARCH64_BTI_PLT and DT_AARCH64_PAC_PLT tags to the dynamic section.

6.4.2 Dynamic Section

AArch64 specifies the following processor-specific dynamic array tags.

Table 6.4: Table 5-4, AArch64 specific dynamic array tags

Name	Value	d_un	Executable	Shared Object
DT_AARCH64_BTI_PLT	0x70000001	d_val	Platform specific	Platform Specific
DT_AARCH64_PAC_PLT	0x70000003	d_val	Platform specific	Platform Specific
DT_AARCH64_VARIANT_PCS	0x70000005	d_val	Platform specific	Platform Specific

DT_AARCH64_BTI_PLT indicates PLTs enabled with Branch Target Identification mechanism.

DT_AARCH64_PAC_PLT indicates PLTs enabled with Pointer Authentication.

The presence of both DT_AARCH64_BTI_PLT and DT_AARCH64_PAC_PLT indicates PLTs enabled with both Branch Target Identification mechanism and Pointer Authentication.

DT_AARCH64_VARIANT_PCS must be present if there are R_<CLS>_JUMP_SLOT relocations that reference symbols marked with the STO_AARCH64_VARIANT_PCS flag set in their st_other field.